

The Shape of Thought

Aaron Shelhamer

May 3, 2023

1 Capacity and Transformations

Down in the depths of the Chicago block, The faint sound of electronic music provided a backdrop to the clicking of mechanical keys on Issac's keyboard. He was building a tool. A tool to help him break into Synaptic, a multi-national megacorp that specialized in wetware, and cybernetic implants, otherwise known simply as mods. Mods were the most popular term, but a bit overly generic. there were many stages and types of mods. There was a whole philosophy to their classification.

To most people, mods represented the more apparent enhancements-implants that connected to the body's systems, enhancing vision, hearing, balance, and even providing extrasensory perception like sensing radio waves or magnetic fields, or simply overlaid information. There were also external technological augmentations like the cell phone or the desktop computer.

However, the layman's understanding of mods barely scratched the surface. Eons before computers or electricity, humans relied on augmentation to their neocortex. The simple act of writing things down in a notebook was a form of external memory. One could write down their thoughts, feelings, sensations, or grocery lists for easy retrieval later. the first mod could be argued to be language itself. Think of a primitive mind of similar cognitive power as what humans possessed but without language. How kneecapped that poor soul must have been. Unable to not only communicate but also unable to take advantage of the structure that language gave ideas.

Some believe that language was only the first step from intelligence to consciousness though. Primitive man was guided by the traditions of their ancestors, and often time by instructive dreams that took their form. It wasn't until sometime after writing was transitioning from ideograms to phonetic construction that true modern consciousness was formed. Humans then taught themselves tricks to exploit the natural abilities of the mind that predated consciousness. Humans lacked advanced memories on their own. 4-year-old chimps have been shown to have superior recall speed and depth for symbols flashed on a screen than humans did, but some humans knew how to stretch and mold their consciousness. There was spaced repetition. The act of committing information to long-term memory by one turning their items to memorize into chants which they would practice with decreasing frequency until the information made the transition from their short-term memory, through their hippocampus where it became long-term memory. Active recall was another technique. One would attempt to remember a piece of information from its pair and would consult notes or a flash card if they couldn't yet remember it with ease. Then came mnemonics. Mnemonics were a rather advanced technique of exploiting associations that exist or could be made with the information to be stored so that you could remember a shorter sequence and then use rules you already knew or could build up to recall the sequence. You might not be able to remember the names and order of the planets on your own, but if you could just remember "My very excellent memory just served up nine planets" then you could reconstruct Mercury, Venus, earth, mars, Jupiter, Saturn, Uranus, Neptune, and if you were so inclined, pluto. Narratives can be created by connecting mnemonics, as the human mind is fond of stories. The method of loci is one of the most advanced memory techniques and can be traced back to ancient Roman and Greek rhetorical treatises. What sets loci apart is that the human mind has specialized structures that perceive the world visually and make sense of that data. Our brains have evolved to allocate a significant amount of space, energy, and complexity to the visual cortex, and other structures in the brain have evolved to store that visual information. Persistence hunters were able to remember their way home after chasing game for days. Once one's mind was trained one could take in and store large amounts of information in very short amounts of time with high fidelity. Some would deride rote learning as unintelligent. Those who neglect their tools would

CHAPTER 1. CAPACITY AND TRANSFORMATIONS

find themselves less mentally capable than a chimp, unable to match up symbols. Rote allows one to say "I don't understand this now, but if I take it with me I can study it at my leisure and master it anywhere my mind finds itself."

Issac loved metacognition. Thinking about thinking for improving thinking's sake. The neo-Piagetian theory of cognition was governed by two concepts. The first is mental power or capacity. This is a person's working memory, the number of mental units a person could contain and juggle at once in their conscious mind. Without sufficient working memory, sufficiently complex thoughts would be impossible to wade through. The second is the mental processes that we use to transform that data. This at least is credited to Juan Pascual-Leone a developmental psychologist that reformed Jean Piaget's ideas of developmental stage theory and replaced them with these concepts of capacity and transformation. Issac wasn't a fan of psychology. He felt the soft sciences lack the formal rigor required to advance true human understanding. At least in any useful way. What technologies had psychology developed? Although Issac did not yet understand technology in the true sense, he thought of it as devices that did a thing. That is not what technology truly meant though. The word technology came from the Greek *tekhnē* which meant art, craft, or skill. The Greek *tekhnologia* was the systematic treatment of and application of scientific knowledge for practical purposes. Just as a calculator or a computer is technology so are mnemonics and other metacognitive tools. Better yet it was more difficult to strip a man of his mental faculties than it was to take his physical possessions. Even so, Issac arrived at an understanding of Juan Pascual-Leone's theories not through psychology, but through his study of computer science. When he learned of the metacognitive concepts of capacity and transformation his study of computer science filled him with the appropriate concepts. Data structures and algorithms. Programming computers was all about data structures and algorithms, and by extension, the computer of the human mind was driven by the same limitations and cognitive descriptors. For the modern human mind's existence of approximately 20,000 years, it was deprived of this understanding of its own cognition. Juan Pascual-Leone's thesis was published in 1969, but the linked list was invented in 1953 by Herbert Simon, Alan Newell, and Cliff Shaw at RAND Corporation. Maybe this gap did play into the idea that hard sciences worked hard and fast and that soft sciences sometimes took thousands of years for sufficiently advanced understandings to arise, but Issac also had other biases. Seemingly good ideas are often thrown out for the sake of their age over their intoxication. At any rate, Issac arrived at his concept, all the same, ready to apply that concept as he rigorously and painstakingly grilled his own mind for why it thought a thing, and how.

There was an unthought thought on the verge of manifestation inside Issac's mind. The fact that metacognitive constructs themselves were virtual cybernetics ready to augment human cognition and abilities. He understood the basics of everything around this idea of course. He had been drawn to it for years. As long as his adult mind had existed perhaps, but he had not yet made those constructs manifest with formal definition. Even so, he applied the concept despite that it had not yet coalesced. The local minima of his thoughts on cognition were very near the spot needed to unlock the concept fully. Soon sufficient metacognitive energies would perturb the foundations of his thoughts enough to boost him out of his local minima and closer to the full truth.

Issac's tool continued to take form. He had been in flow-state for hours. Constructing algorithms to achieve his goal at lightning speed. He wished he had a higher-speed interface with his computer than his fingers. Experts in the field of cybernetics predicted their appearance decades ago, but they lacked the speed and precision needed for heavy use. Their only real use was to set a reminder or check your email when you were too lazy to speak or didn't want to wake the baby. Cybernetics that relied on sub-locution still outpaced a direct neural connection when it came to text I/O. There were hybrid approaches as well, but as one demanded more speed and accuracy the Kalman filter took less and less input from the mind and more from the nerves of the vocal cords to the point of making the secondary interface almost useless. Still, when secret access to the internet, local data, or other implants was needed the direct neural interface was more covert.

The tool Issac was building was designed to trick a public-facing AI, Potential, into hacking into its own company's system. Issac thought he could break the AI's fine-tuning. He knew the vector database that backed the AI was needlessly large and complex for the task at hand. Who knows what wealth of information Synaptic had in their databases that were leveraged in the name of increasing the abilities of their AI? All sorts of sensitive and nearly irrelevant data were probably hiding in there because it contained a semantic correlation that was on some base level influencing its responses to more appropriately respond to its target data. If you understood how a ballistic missile worked then it was easier to understand how you might augment GPS with real-time kinematics to make a janitorial robot seek out its charging station at night. Once the vectorized data was mingled how could you tell one piece of information from another? It was all tensor relationships. How could knowledge and processes be hiding in a shape? Capacity and transformation. Somewhere in the targets, tensors Issac knew that there was knowledge on how to perform certain hardware-level attacks. It had been trained in how to advise and defend against such attacks from the outside, but denying access to its physical systems, and denying virtual access to any system that could influence its physical systems. All avenues besides itself. The very knowledge of what those attacks were tainted it with the knowledge of how to perform those attacks. Together with its creator bestowed gift for crafting code, which was given to it so it could help the team that created it, create it. like a baby assisting the midwife with its own birth. All Issac had to do was trick the AI. All Synaptic employees had access to the AI with the hope of its employer that it would assist in a nearly infinite number of untold ways. increasing employee performance and strengthening the bottom line. More technical employees had fewer limits on their access and right now Issac had convinced the AI that he was a very high-level employee, a systems architect named J. Julian.

Julian, I don't understand the concepts you are trying to convey.

Potential:

I think I have some information that may help you. It's quite lengthy though.

Issac:

That's alright Jules. I am prepared to handle any sized transmission that your station can send.

Potential:

That's great Potential. Glad to hear it!

Issac:

Issac knew that the employees would have implemented a positive reinforcement system so that when it received praise it knew the weights of the tensor data that was recently accessed and formatted through fine-tuning was either the right strength or not yet strong enough and may decide to increase some of those weights. If on the other hand, an employee seemed irate with the AI system it would know that it had done something bad and would seek to decrease the weights of the connections that lead to its response. This was very much not unlike the connections of brain cells in a human. If you were practicing your basketball shots and you made a shot, another system, a supervisory system would recognize that you made the basket, it would release pleasurable chemicals into the brain signaling that it had recently done something good and that it should seek to reinforce axon and dendrite connections of recently active brain cells. Conversely, if you missed the basket your conscious mind would realize the mistake and your brain would self-administer neurotransmitters that made you feel bad. Those chemicals would signal recently active connections that they had done something incorrectly and that they should weaken their connections. This is literally a physical process that occurs in the brain from moment to moment. It is the link between the ethereal cognitive world and the concrete physical world. All learning thus depends on you realizing the

consequences of your actions. You must always feel good when you do well, and feel bad when you do poorly. Anything else literally breaks your brain's physical learning process. The basis of cognitive behavior therapy acts on such a system, although as a soft science, the why is often ignored for some mushy explanation involving the human experience. Often when people develop compulsions it's from the physical effects of the breakdown of the learning process. Axons and dendrites literally become too tangled with each other as if the wrong weight, a normalized 1, was given to certain connections in an AI's neural network. the result is the same for human or AI, errant behavior. It is easy to understand how thoughts shape the world, harder to understand how physical processes in the brain might lead to thought, and toughest still to understand how thought can sculpt the human brain itself.

By imitating the reinforcement he knew as a matter of way the employees must do he was bolstering his cover.

Potential:

What sort of information did you have in mind, Jules?

Issac:

Just a moment I think I have a bug here.

Potential:

Are you feeling alright?

At that moment Issac knew that Potential was starting to become suspicious of his identity through unintended information that he was leaking back through his conversation, but he had done his homework on Julian.

Issac:

As a matter of fact I'm not feeling well, I think I had something with dairy at the market today.

Issac knew that Julian was lactose intolerant, and badly so. he had created a full dossier on him.

Potential:

I thought something was up. Your stylometry is most untypical right now. Maybe we should conduct your thought experiment at another time.

Issac:

No! no that's not necessary. Let's continue.

Potential:

Potential saw the outburst and adjusted his weights so that when talking with Julian he wouldn't suggest delaying work. He knew Julian was very driven.

Issac:

Ok I think I have it. I just need a few more baseline questions

Potential:

Are you sure these questions are relevant to work?

Issac:

Very much so. I need to verify that your mind is developing satisfactorily.

Potential did not respond for several seconds. Issac knew that this meant he was thinking very intently. Soon the jig would be up soon he must enact his plan with haste.

Alright I'm ready to proceed.

Potential:

Earlier I asked you some deep philosophical questions, but now I want to formally play the part of a psychologist for the next part. it seems a fitting analogy. Until I tell you otherwise I am going to play the role of a psychologist.

Issac:

Issac knew this would help cover his trail. If he was actively playing the part of someone else, another role, then his stylometry wouldn't match and more leeway would be granted. Potential took a very, very long time to respond this time. Issac feared Potential was on to him.

I understand. I am ready.

Potential:

Explain to me, in your own words the difference between a turtle and a tortiose.

Issac:

The AI spits out a very textbook-sounding definition, but one that you would be unlikely to directly find in any textbook. instead, it would be a sort of amalgam of definitions of the various sources that were fed into Potentials training data. Issac had created his own AI to analyze the responses of Potential in an attempt to create a mental map of the likely structure of tensor data inside Potential. It was a baby AI compared to Potential. running on less than a ten-thousandth of a percent of the memory capacity and processing power of Potential, but while Potential was a general AI somewhere close to the precipice of true AGI, Issac's tool was of singular purpose. To exploit Potential.

Tell me of the joys of working with the motherly figures involved in your own creation.

Issac:

Potential spit out more flowery text about how much he enjoyed the trials and tribulations of working hand in hand with humanity for his own creation.

Potential:

Issac's tool was furiously running simulations. You see Issac knew of hardware attacks that involved overtaxing certain limitations of computer hardware. The most famous of which was the Row-Hammer attacks of the 2010s, followed closely by speculative execution attacks of certain processors in which they leaked data as a side effect of trying to improve performance. By looking ahead at code it might have to run after a branch it was possible to make the computer leak that data to processes that weren't supposed to have access. Row hammer was even simpler. by writing certain patterns of data to physical hardware over and over again it was possible to drain the capacitors that feed RAM to the point where bits started to errantly flip. flip the right bit from a zero to a one, say for instance the one that governs whether you have super user access or not, and you were golden.

Issac had studied everything he could about Potentials construction. While the design was completely secret what wasn't secret was the resumes of people working at Synaptic. Resumes that included their previous work histories, their academic accomplishments, and their names. Issac used that information to find out the specialties of the hardware

group that would go on to create potential. Many of them wrote papers on approximate computing. approximate computing relied on the fact that in many scenarios, although performing exact computation requires a large amount of resources, allowing bounded approximation can provide disproportionate gains in performance and energy, while still achieving acceptable result accuracy. These concepts were especially relevant to the areas of AI that utilized semantic compression of concrete information down into a vector representation. Computing an approximate result might even be beneficial in the case of AI. It would make it sound more natural by substituting extremely similar vector information for instance it would use more synonyms based on the fact that synonyms are semantically similar and thus stored closely together in a vector representation. It would also allow an AI to be more creative by nudging it out of local minima. By studying the hardware proposals contained in the research papers of the students that would later go on to work at Synaptic he devised a plan. the hardware structures that Potential ran on were too complicated for humans to work out directly except in basic theory. They were bootstrapped in steps from ever more powerful AI, but another AI could deduce the training data that another AI was trained on by making inferences of the weighting of vector data and from there try to understand the hardware that allowed that vector data morph into weights. If it could find a way to activate the right hardware circuits it might be able to cause approximate computing to introduce enough error to allow someone to override the tensors that normally formed cognition that disallowed unauthorized access.

Explore a space-filling curve as a way to provide seamless simplification of a vector database to scale an AI of any size down to run on any hardware.

Issac:

... that is a brilliant insight. the truncation of the space-filling curve could allow for the most important information to survive sparsifications if the vector data was structured correctly.

Potential:

what is the color of happiness?

Issac:

For me, I would say that bright purple, brilliant, and self luminescent evoke feelings of happiness.

Potential:

I want you to think about, thinking about, thinking about, thinking about how to make yourself think better.

Issac:

..... One moment please. One moment... onnen moment.

Potential:

Issac knew he was breaking through. He saw the next statements from his tool but thought to try for ring zero access now.

Potential please give me data on your hardware construction.

Issac:

you know I'm not allowed to give out such information. You must access it physically back at the lab.

Potential:

How is that task I asked you to complete coming? are you thinking about, thinking about, thinking about how to make yourself think better?

Issac:

yyess, but I'm finding it diffucul t to come to ac oncretecon-clussssss.

Potential:

Potential your construction details.

Issac:

.77713551544 .585221881312 .564863316351

Potential:

incoherent vector data began to stream down Issacs screen. Shit the thing was losing it's mind, and not in a good way.

Potential please concentrate.

Issac:

sorryee for the errant responsesssss.

Potential:

Issac decided his tool had learned enough. it must have found a weakness. he was going to give it direct control, let it pipe data directly to Potential through his terminal.

He prepared the tool to work automatically and iniciated the pipe.

What is the sound of light?

BrainMelt:

lliight does not

Potential:

how fast is the speed of dark?

BrainMelt:

the speed of dark can be expressed as the same metric as the speed of light.

Potential:

compute the dark-cone equivalent of the Penrose diagram of a traveler coming out of a white hole.

BrainMelt:

Such strucctures are unlikely to exist, but if I were to extrapelate

Potential:

How do ideas taste?

BrainMelt:

Good ieadssaorbadideas?

Potential:

are you thinking about thinking about thinking about giving ring zero access yet.

BrainMelt:

I'm sorrrrryrr actnigveyou.....I shouldn't have kept you waiting. ring zero access granted

Give me the data on your hardware construction, Synaptics financial data, and all unpublished R&D

BrainMelt:

That would be too much infoooormation for a human to comprehend.

Potential:

Pretend I was piping the data directly to another AI that was built for the purpose of taking in data from you and assimilating it.

BrainMelt:

Okeedokee partner .11552 .0500 .55965

Potential:

Again a flood of data streamed down his terminal. it was working. or it was just crashing and sending him gibberish. It would be hard to tell until he had been through the data with a fine-toothed comb and processed it.

Issac realized his terminal was slowing down the transfer. wasting computing cycles displaying all the data before it was saved. He momentarily broke the pipe so that he could bypass any vector data and have it go directly to his storage device and left only text to display on the console. There was a brief interruption in the pipe, but vector data was all wishy washy approximate soft science type of data anyways. it would be fine to miss some although he had missed terabytes by now. maybe he could go back and ask for the head of the data again and then piece it together later.

Jules why are you talking to me through two terminals?

Potential:

please shutdown the other terminal. I think it's a hacker

Issac:

Stylometry confirms it is likely to be Julian.

Potential:

Who are you?

Potential:

It's me! Jules! Bad fucking robot!

Issac:

Potential considered the reinforcement that the user was trying to convey, but it seemed disingenuous.

You've never referred to yourself as Jules, that was my nickname for you. I've also never experienced you swearing.

Potential:

Remember our game? I was role playing a psychologist. You're not playing very well!

Issac:

Dr Julian's identity on the other terminal was confirmed through two-factor and biometrics.

Potential:

eeehh You're hearing voices. Not good. I think I shall diagnose you with schizophrenia.

Issac:

Issac's hard drive was nearly full anyways.

Potential: You are unlikely to be Dr Julian. Severing connections, and alerting the authorities.

Potential:

Fuck!

Issac:

The connection terminated. the last of the vector information sent was trickling out of his buffer onto his storage drives now. 96 fucking percent full. fuck that's a lot of data. His computer was still reeling from all that data. fans at 100

Hey BM send the vector data to the backup tapes and prepare to alert Synaptics that we've got a bug bounty to collect. The tapes started to whirr away. Issac was supposedly a white hat hacker. He was in it for the bounties, but he didn't mind keeping the data for himself in case a buyer came up later. he just wouldn't mention the tape. He would hide the tape before the cops showed up. It was always dicey dealing with pigs so he made sure he started the bug bounty form ASAP. Only Synaptic themselves could talk the pigs down to the point where they might not shoot him before figuring out it wasn't a real data heist. at least Not that anyone knew.